

RCNN Architecture Assignment



1. What are the objectives of using Selective Search in R-CNN?
2. Explain the following phases involved in R-CNN:
 - a. Region proposal
 - b. Warping and Resizing
 - c. Pre trained CNN architecture
 - d. Pre Trained SVM models
 - e. Clean up
 - f. Implementation of bounding box
3. What are the possible pre trained CNNs we can use in Pre trained CNN architecture?
4. How is SVM implemented in the R-CNN framework?
5. How does Non-maximum Suppression work?
6. How Fast R-CNN is better than R-CNN?
7. Using mathematical intuition, explain ROI pooling in Fast R-CNN.
8. Explain the following processes:
 - a. ROI Projection
 - b. ROI pooling
9. In comparison with R-CNN, why did the object classifier activation function change in Fast R-CNN?
10. What major changes in Faster R-CNN compared to Fast R-CNN?
11. Explain the concept of Anchor box.
12. Implement Faster R-CNN using 2017 COCO dataset (link: <https://cocodataset.org/#download>) i.e. Train dataset, Val dataset and Test dataset. You can use a pre-trained backbone network like ResNet or VGG for feature extraction. For reference implement the following steps:
 - a. Dataset Preparation:
 - i. Download and preprocess the COCO dataset, including the annotations and images.
 - ii. Split the dataset into training and validation sets.
 - b. Model Architecture:
 - i. Build a Faster R-CNN model architecture using a pre-trained backbone (e.g., ResNet-50) for feature extraction.
 - ii. Customise the RPN (Region Proposal Network) and RCNN (Region-based Convolutional Neural Network) heads as necessary.
 - c. Training:
 - i. Train the Faster R-CNN model on the training dataset.
 - ii. Implement a loss function that combines classification and regression losses.
 - iii. Utilise data augmentation techniques such as random cropping, flipping, and scaling to improve model robustness.
 - d. Validation:
 - i. Evaluate the trained model on the validation dataset.
 - ii. Calculate and report evaluation metrics such as mAP (mean Average Precision) for object detection.
 - e. Inference:
 - i. Implement an inference pipeline to perform object detection on new images.
 - ii. Visualise the detected objects and their bounding boxes on test images.
 - f. Optional Enhancements:
 - i. Implement techniques like non-maximum suppression (NMS) to filter duplicate detections.
 - ii. Fine-tune the model or experiment with different backbone networks to improve performance.

Submission Guidelines:

- Answer all the questions in a single Jupyter Notebook file (.ipynb).
- Include necessary code, comments, and explanations to support your answers and implementation.
- Ensure the notebook runs without errors and is well-organized.
- Create a GitHub repository to host your assignment files.
- Rename the Jupyter Notebook file using the format "date_month_topic.ipynb" (e.g., "21st_September_RCNN Architecture.ipynb").
- Place the Jupyter Notebook file in the repository.
- Commit and push any additional files or resources required to run your code (if applicable) to the repository.
- Ensure the repository is publicly accessible.
- Submit the link to your GitHub repository as the assignment submission.

Grading Criteria:

- Understanding and completeness of answers: 40%
- Clarity and depth of explanations: 25%
- Correct implementation and evaluation of optimizer techniques: 15%
- Analysis and comparison of different optimizers: 10%
- Proper code implementation and organization: 10%

Note: Create your assignment in Jupyter notebook and upload it to GitHub & share that uploaded assignment file link through your dashboard. Make sure the repository is public.