



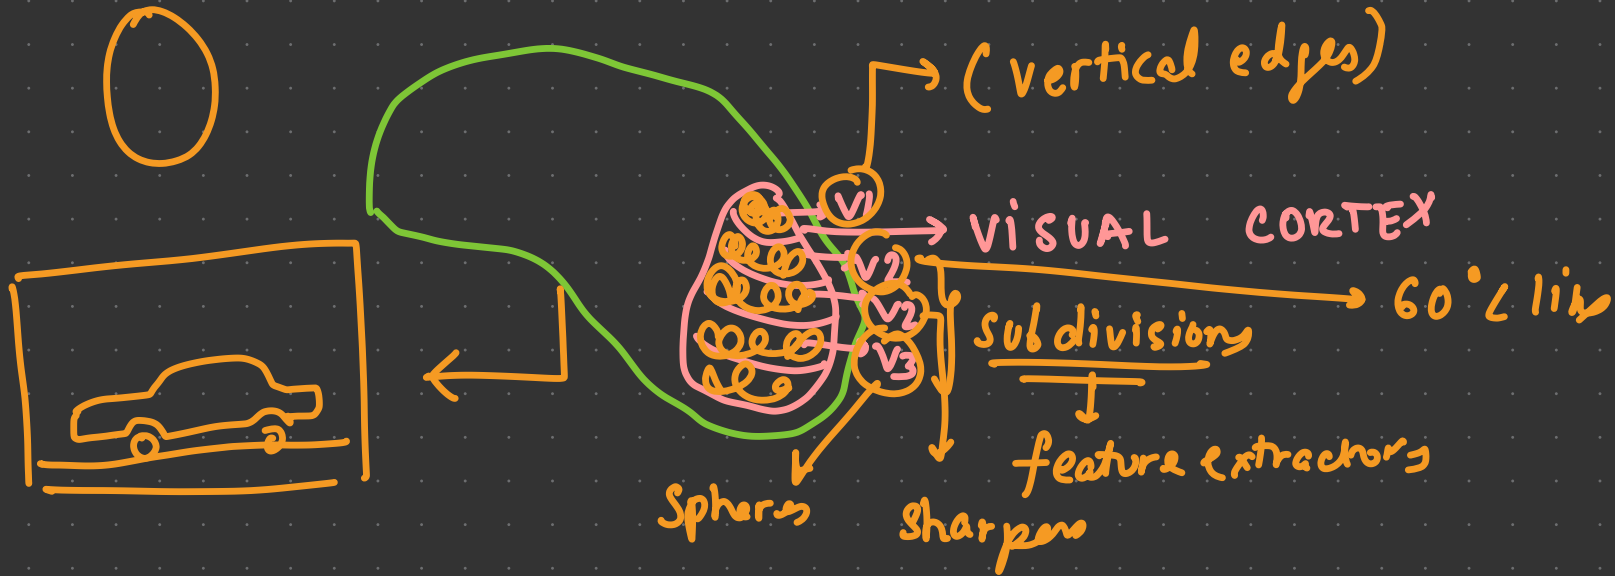
CONVOLUTION NEURAL NETWORKS


1 →

USE


→ these are used to work on
visual data (videos & images)

How does a person recognize different objects





The CNN Complete Network Overview



Im \rightarrow Conv + MP \rightarrow Conv + MP

IMAGE

GRAY SCALE


RGB (MULTIPLE CHANNELS)

0-255


$n \times n$

10x10

w




GRAY





GRAY CHANNEL

RGB



★ CONVOLUTION]





15000×1000

0
0
0
0
0
0
0
0
:
:
:

10 Labs \rightarrow 1 Billion

(Weights & biases)



CONVOLUTION

$$\begin{array}{c|c} 0.1 & 0 \\ \hline 0 & 0.1 \end{array}$$

Input

6x6

3	2	2	1	2	4
0	1	1	2	1	0
1	1	2	2	0	1
2	0	0	1	0	2
1	2	1	0	2	2
0	2	1	0	1	3

0.3 + 0.1 (Feature Map)

0.4	0.3	0.4	0.2	0.2
0.1	0.2	0.3	0.2	0.2
0.1	0.2	0.2	0.2	0.2

FM

5x5

(FILTERS)

kernel

F1

0.1	0
0	0.1

vertical edges

F2

1	5	1
2	1	5

horizontal edges

F3

2	2
2	1




sharpen

0.4	0.1	0	0.3	0.2
0.3	0.3	0.1	0.1	0.5

→ We were able to extract information from your image using a filter

100 Feature Maps

* FILTER → ?
 * DIMENSIONS → ? → (PADDING) × (STRIDES)
 (PADDING)



Step Size = 1

→ STRIDE

✓ → USE CASES CNN

✓ → HUMAN BRAIN → DETECT IMAGES
and EXTRACT FEAT

✓ → CONVOLUTION

✓ → PADDING

✓ → STRIDES


✓ → FILTERS

✓ → POOLING


} DIMENSIONS 128D

0.5	0.7
0.3	0.4




FEAT. EXTRACTION



WEIGHTS




feature map



28x28

14x14



POOLING

stride

max
min
avg

pooling
pooling
pool.

(2x2)

4	3	5	1
9	2	7	4
2	3	6	5
2	0	4	8

Minimum value
in
min pool.

9	7	9
9	7	7
3	6	8

Avg pooling → avg.

IMPORTANT FORMULAS

Input Dimensions:

Width (W): The width of the input image.

Height (H): The height of the input image.

Channels (C): The number of channels or color planes in the input image (e.g., 1 for grayscale, 3 for RGB).

Convolutional Layer:

Kernel Size (K): The width and height of the convolutional kernel or filter.

Padding (P): The number of zeros added to the border of the input image to preserve spatial dimensions.

Stride (S): The step size at which the kernel moves across the input image.

Number of Filters (N): The number of convolutional filters applied in the layer.

Output Dimensions:

→ Width (W_{out}): $((W - K + 2P) / S) + 1$

→ Height (H_{out}): $((H - K + 2P) / S) + 1$

→ Channels (C_{out}): N

Pooling Layer:

Pool Size (P): The width and height of the pooling window.


Stride (S): The step size at which the pooling window moves across the input.



Output Dimensions:

Width (W_{out}): $((W - P) / S) + 1$

Height (H_{out}): $((H - P) / S) + 1$

Channels (C_{out}): C (remains the same)





5 0 4 1 9 3 2

0 9

→ [0 0 0 0 0 1 0 0 0]

→ [1 0 0 0 0 0 0 0 0]

→ [0 0 0 0 0 1 0 0 0]

